

Model drift is a phenomenon in real-world machine learning (ML) deployments where model performance can degrade over time due to shifts in data distribution. A model monitoring pipeline can be implemented to ensure model reliability and maintain high predictive accuracy.

A robust model monitoring pipeline consists of several key components. The following sections outline each component and suggest suitable tools to use.

### **Establishing Baseline Model Performance**

A baseline serves as a reference for detecting deviations in production data. Baseline performance is established using validation metrics such as accuracy, precision-recall, F1-score, and AUC-ROC. Additionally, a distribution profile of input features and predicted outputs should be recorded.

### **Data Ingestion and Logging**

This captures incoming feature values, predictions, and ground-truth values where available. This can be achieved using logging frameworks such as WhyLogs, MLflow, or cloud solutions such as AWS SageMaker Model Monitor.

### **Drift Detection and Statistical Analysis**

With the data collected, the system should periodically compare production data against the baseline using statistical drift detection techniques. Common methods include: Population Stability Index (PSI), which measures shifts in feature distributions over time; Kullback-Leibler (KL) Divergence, which quantifies how much the current feature distribution differs from the training data; Jensen-Shannon Divergence, a symmetric variation of KL divergence used for more stable drift detection; and Mean and Standard Deviation Shift, which tracks numerical feature variations to detect anomalies. Automated tools such as Evidently AI and Great Expectations can be integrated into the pipeline to monitor these statistical properties efficiently.

### **Model Performance Evaluation**

The above sections covered feature drift detection and while that is important, tracking model performance metrics is crucial. If the model's accuracy, precision, or recall declines beyond a predefined threshold, this may indicate **concept drift**, requiring model retraining. Periodic evaluations should be performed using real-world labeled data where available.

### **Alerting and Decision Automation**

When acceptable thresholds are exceeded, the pipeline should incorporate automated alerting mechanisms to notify the team in-charge. Monitoring platforms such as Prometheus, Grafana, or Weights and Biases can do this.

## **Model Retraining and Deployment**

Once drift is detected, the pipeline should automatically trigger a model retraining workflow using updated data. Before deployment the retrained model must be validated against benchmark performance metrics to ensure improved or stable performance. Modern ML platforms such as MLflow, Kubeflow, and AWS SageMaker, enable automated retraining and version-controlled deployments.

## **Integration with GitHub Actions for Automation**

There are many combinations of open-source and enterprise-grade tools available for implementing continuous monitoring; these have been mentioned above. To orchestrate these steps seamlessly, GitHub Actions can serve as an automation engine. A GitHub Actions pipeline can trigger different stages of the model monitoring pipeline, such as:

- Running scheduled drift detection scripts.
- Retraining the model when drift is detected.
- Deploying the updated model to production.
- Logging model versions and monitoring data distribution.

By integrating GitHub Actions with various tools, organizations can fully automate model monitoring, ensuring models remain accurate and robust in production.